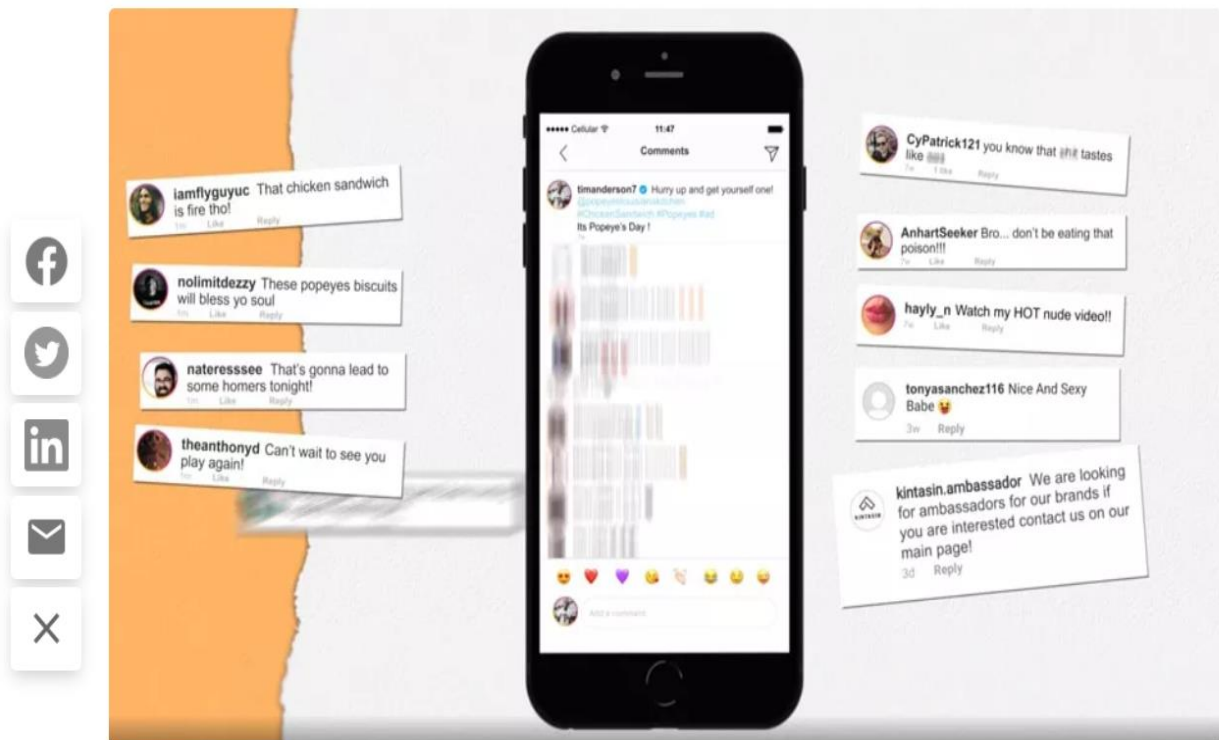




Malignant comment Classifier Project





Submitted by:

SARMISHTHA HALDAR

ACKNOWLEDGEMENT:

I would like to thank all my teachers, supervisors for the learning especially Shubham Yadav Sir. Few journals referred in the case are as follows:

i) Yin Dawei ,et al *Detection of Harassment on web 2.0*

ii) Liang-Chih Yul Jin *Refining word embeddings for Sentiment Analysis.*

Business Problem Framing

This is a classic problem in today's world of social media platform where our goal is to build a prototype which helps to evaluate the hate and abusive comments made on social media platform so that it can be controlled and restricted from spreading

Conceptual Background

Online platforms when used by normal people can only be comfortably used by them only when they feel that they can express themselves freely and without any reluctance. If they come across any kind of a malignant or toxic type of a reply which can also be a threat or an insult or any kind of harassment which makes them uncomfortable, they might defer to use the social media platform in future. Thus, it becomes extremely essential for any organization or community to have an automated system which can efficiently identify and keep a track of all such comments and thus take any respective action for it, such as reporting or blocking the same to prevent any such kind of issues in the future. This is a huge concern as in this world, there are 7.7 billion people, and, out of these 7.7 billion, more than 3.5 billion people use some or the other form of online social media. Which means that every one-in-three people uses social media platform. This problem thus can be eliminated as it falls under the category of Natural Language Processing. In this, we try to recognize the intention of the speaker by building a model that's capable of detecting different types of toxicity like threats, obscenity, insults, and identity-based hate. Moreover, it is crucial to handle any such kind of nuisance, to make a more user-friendly experience, only after which people can actually enjoy in participating in discussions with regard to online conversation

INTRODUCTION

Internet comments are bastions of hatred and vitriol. While online anonymity has provided a new outlet for aggression and hate speech, machine learning can be used to fight it. The problem we sought to solve was the tagging of internet comments that are aggressive towards other users. Here our goal is to build a prototype of online hate and abuse comment classifier which can be used to classify hate and offensive comments so that it can be controlled.

Problem Statement

To Build a model which can be used to classify hate and offensive comments.

Analytical Problem Framing

We have used methods like Accuracy, Confusion matrix and Roc-auc curve for model evaluations

Accuracy is defined as the percentage of correct predictions

A **Confusion matrix** is an $N \times N$ matrix used for evaluating the performance of a classification model, where N is the number of target classes. The matrix compares the actual target values with those predicted by the machine learning model. This gives us a holistic view of how well our classification model is performing and what kinds of errors it is making.

		ACTUAL VALUES	
		POSITIVE	NEGATIVE
PREDICTED VALUES	POSITIVE	TP	FP
	NEGATIVE	FN	TN

Let's decipher the matrix:

- The target variable has two values: **Positive** or **Negative**
- The **columns** represent the **actual values** of the target variable
- The **rows** represent the **predicted values** of the target variable

ROC Curve :The Receiver Operator Characteristic (ROC) curve is an evaluation metric for binary classification problems

- Data Sources and their formats

We received the data in the form of .csv file and data was loaded using Pandas

```
In [4]: train=pd.read_csv("Malignantcomment-train.csv")
```

```
In [6]: test=pd.read_csv("Malignantcomment-test.csv")
```

- Exploratory Data Analysis

In [5]: `train.head()`

Out[5]:

	id	comment_text	malignant	highly_malignant	rude	threat	abuse	loathe
0	0000997932d777bf	Explanation\nWhy the edits made under my usern...	0	0	0	0	0	0
1	000103f0d9cfb60f	D'aww! He matches this background colour I'm s...	0	0	0	0	0	0
2	000113f07ec002fd	Hey man, I'm really not trying to edit war. It...	0	0	0	0	0	0
3	0001b41b1c6bb37e	"\nMore\nI can't make any real suggestions on ...	0	0	0	0	0	0
4	0001d958c54c6e35	You, sir, are my hero. Any chance you remember...	0	0	0	0	0	0

In [7]: `test.head()`

Out[7]:

	id	comment_text
0	00001cee341fdb12	Yo bitch Ja Rule is more succesful then you'll...
1	0000247867823ef7	== From RfC == \n\n The title is fine as it is...
2	00013b17ad220c46	" \n\n == Sources == \n\n * Zawe Ashton on Lap...
3	00017563c3f7919a	:If you have a look back at the source, the in...
4	00017695ad8997eb	I don't anonymously edit articles at all.

In [9]: `print('train shape is ',train.shape)`
`print('test shape is ',test.shape)`

train shape is (159571, 8)
test shape is (153164, 2)

Descriptive Statistics

```
In [12]: print('train data Set descriptin',train.describe())
print('test data Set descriptin',test.describe())
```

train data Set descriptin		malignant	highly_malignant	rude	threat \
count	159571.000000	159571.000000	159571.000000	159571.000000	
mean	0.095844	0.009996	0.052948	0.002996	
std	0.294379	0.099477	0.223931	0.054650	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	
75%	0.000000	0.000000	0.000000	0.000000	
max	1.000000	1.000000	1.000000	1.000000	

abuse		loathe
count	159571.000000	159571.000000
mean	0.049364	0.008805
std	0.216627	0.093420
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	1.000000	1.000000

test data Set descriptin		id	comment_text
count	153164		153164
unique	153164		153164
top	8955f949a0319327	== GonzosNoze user page == \n\n Why would you ...	
freq	1		1

```
In [12]: print('train data Set descriptin',train.describe())
print('test data Set descriptin',test.describe())
```

train data Set descriptin		malignant	highly_malignant	rude	threat \
count	159571.000000	159571.000000	159571.000000	159571.000000	
mean	0.095844	0.009996	0.052948	0.002996	
std	0.294379	0.099477	0.223931	0.054650	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	0.000000	
50%	0.000000	0.000000	0.000000	0.000000	
75%	0.000000	0.000000	0.000000	0.000000	
max	1.000000	1.000000	1.000000	1.000000	

abuse		loathe
count	159571.000000	159571.000000
mean	0.049364	0.008805
std	0.216627	0.093420
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	1.000000	1.000000

test data Set descriptin		id	comment_text
count	153164		153164
unique	153164		153164
top	8955f949a0319327	== GonzosNoze user page == \n\n Why would you ...	
freq	1		1

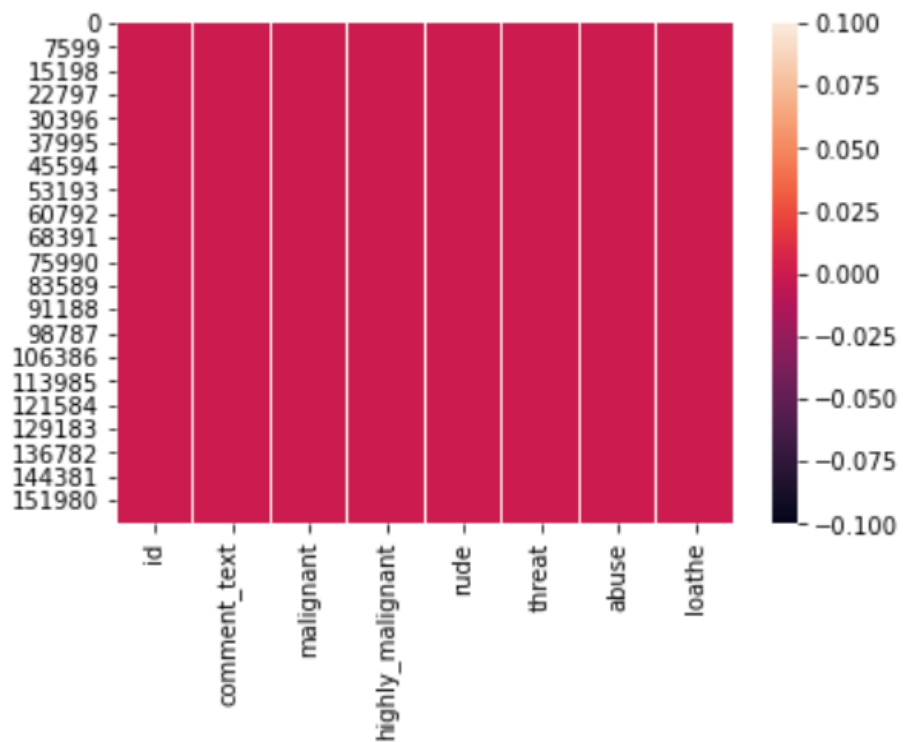
• Data Preprocessing

i)Checking the missing values

```
In [13]: #Checking for missing values
```

```
In [14]: print(train.isnull().sum())  
print(sns.heatmap(train.isnull()))
```

```
id          0  
comment_text 0  
malignant    0  
highly_malignant 0  
rude         0  
threat       0  
abuse        0  
loathe       0  
dtype: int64  
AxesSubplot(0.125,0.125;0.62x0.755)
```



ii) Checking Correlation

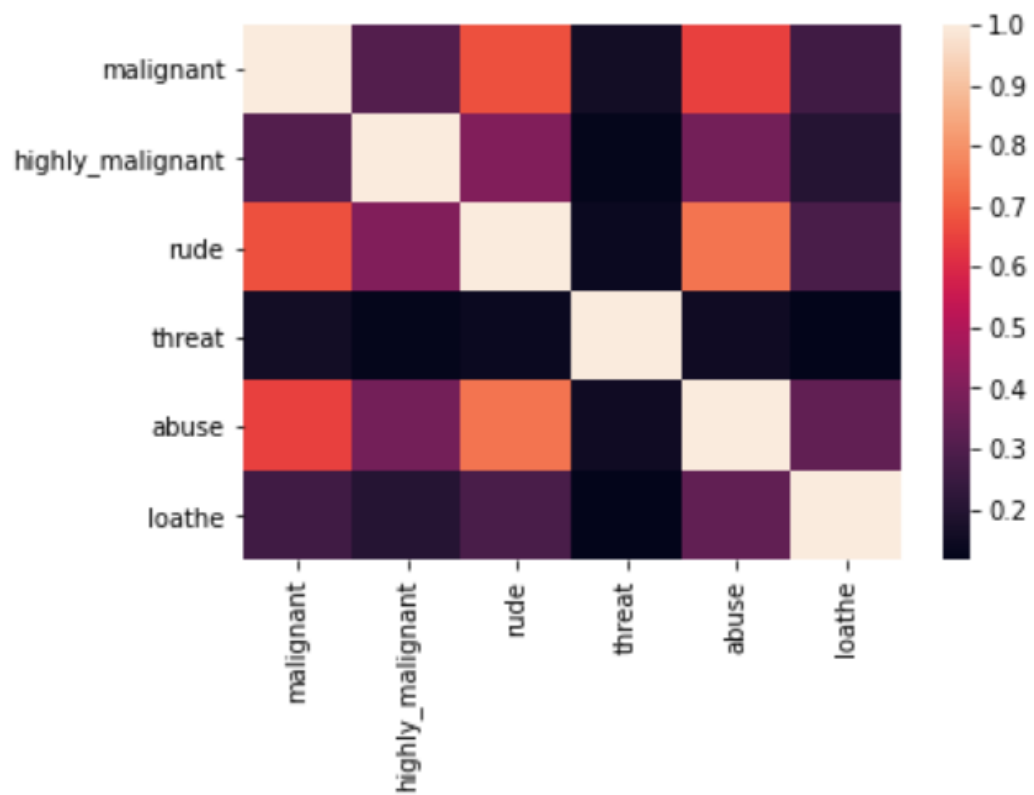
```
In [15]: #Checking correlation in dataset
```

```
print(train.corr())  
print(sns.heatmap(train.corr()))
```


	malignant	highly_malignant	rude	threat	abuse	\
malignant	1.000000	0.308619	0.676515	0.157058	0.647518	
highly_malignant	0.308619	1.000000	0.403014	0.123601	0.375807	
rude	0.676515	0.403014	1.000000	0.141179	0.741272	
threat	0.157058	0.123601	0.141179	1.000000	0.150022	
abuse	0.647518	0.375807	0.741272	0.150022	1.000000	
loathe	0.266009	0.201600	0.286867	0.115128	0.337736	

	loathe
malignant	0.266009
highly_malignant	0.201600
rude	0.286867
threat	0.115128
abuse	0.337736
loathe	1.000000

AxesSubplot(0.125,0.125;0.62x0.755)



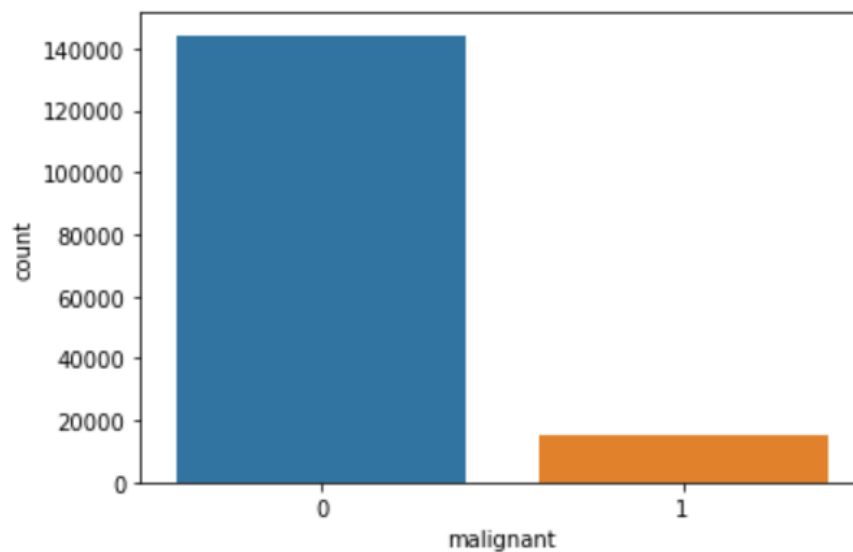
```
In [16]: # checking the skewness for the features:
train.skew()
```

```
Out[16]: malignant          2.745854
highly_malignant          9.851722
rude                      3.992817
threat                   18.189001
abuse                    4.160540
loathe                   10.515923
dtype: float64
```

```
In [17]: col=['malignant','highly_malignant','loathe','rude','abuse','threat']
for i in col:
    print(i)
    print("\n")
    print(train[i].value_counts())
    sns.countplot(train[i])
    plt.show()
```

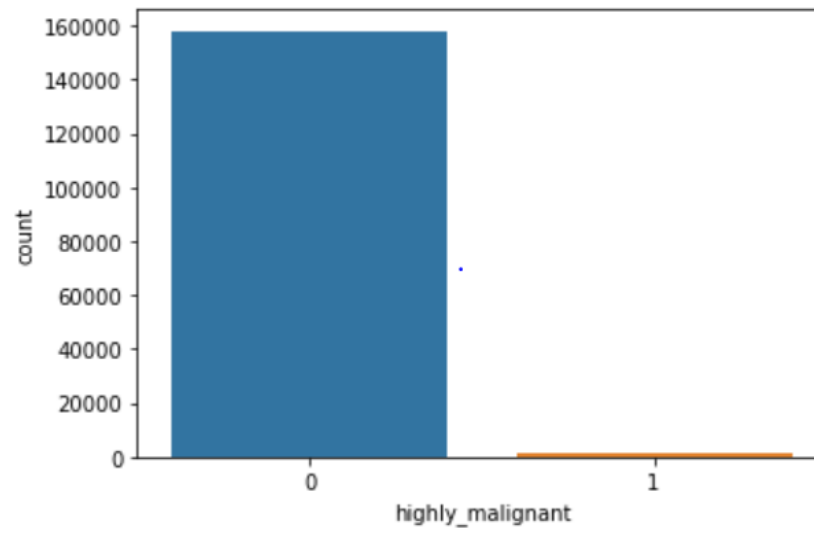
malignant

```
0    144277
1     15294
Name: malignant, dtype: int64
```



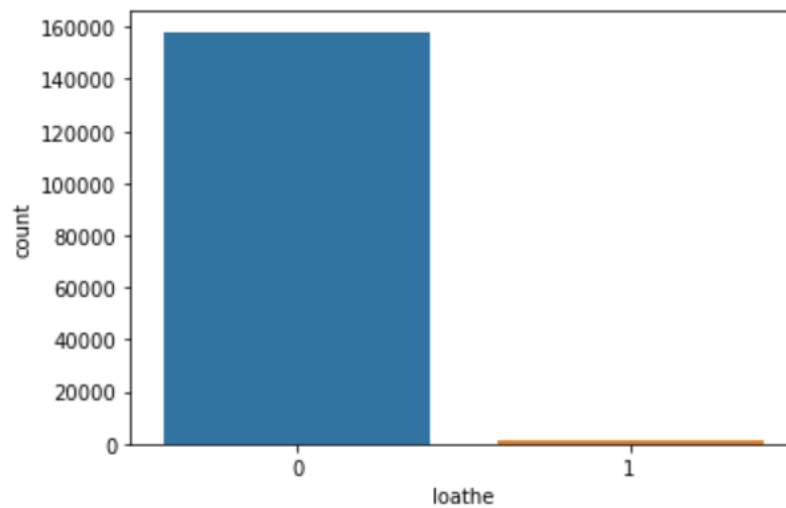
highly_malignant

```
0    157976
1      1595
Name: highly_malignant, dtype: int64
```



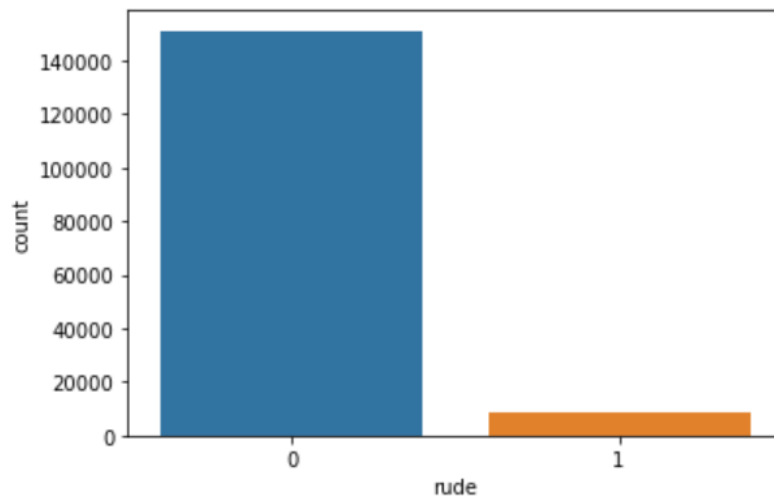
loathe

```
0    158166
1     1405
Name: loathe, dtype: int64
```



rude

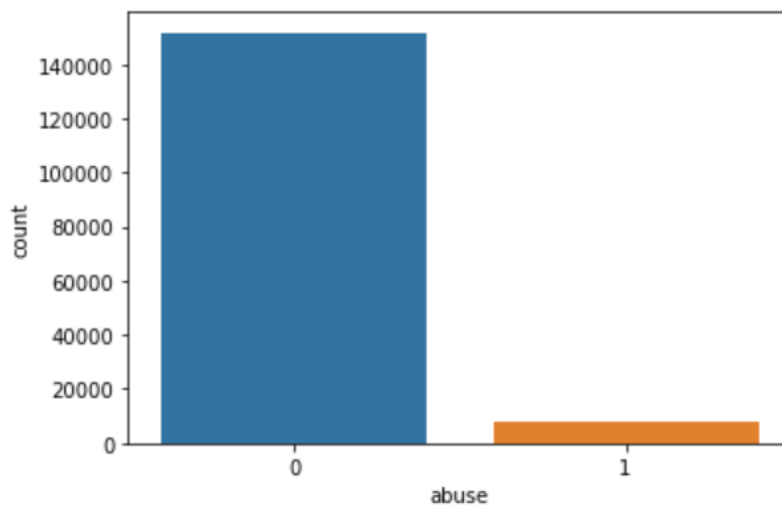
```
0    151122
1      8449
Name: rude, dtype: int64
```



abuse

abuse

```
0    151694
1      7877
Name: abuse, dtype: int64
```



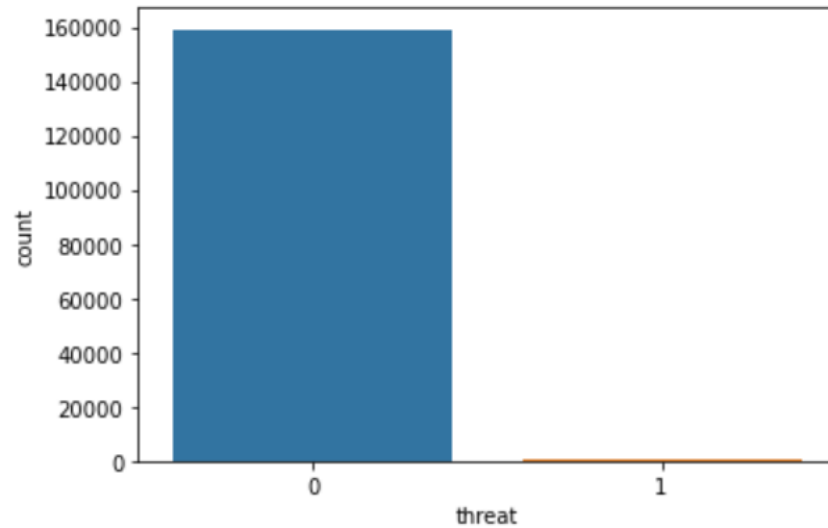
..

threat

0 159093

1 478

Name: threat, dtype: int64



iii) NLP Processing steps carried out to
convert messages to lower case,
replace email addresses with email,
replace URL's with webaddress,
Replace 10 digit phone numbers (formats include paranthesis, spaces, no spaces, dashes) with 'phonenumber',
Replace numbers with 'numbr',
Stop words


```

Origian Length 62893130
Clean Length 43575187

```

[illegible]

```
In [28]: # Convert text into vectors using TF-IDF
from sklearn.feature_extraction.text import TfidfVectorizer
tf_vec = TfidfVectorizer(max_features = 10000, stop_words='english')
features = tf_vec.fit_transform(train['comment_text'])
x = features
```

- **Hardware and Software Requirements , Tools Used**
No Specific requirements except Jupyter Notebook.

Model/s Development and Evaluation

- Identification of possible problem

This is a Classification problem and we have used Decision tree Classifier, Random Forest Classifier, XGboost, Adaboost, K-Neighbors classifier to build the model as per the performance the Random forest classifier is chosen to be the best one.

- Run and Evaluate selected models: We ran and evaluated the above mentioned models and chose Random forest classifier as the best model.

- **LOGISTIC REGRESSION:**

Logistic regression is a statistical model that in its basic form uses a logistic function to model a binary dependent variable, although many more complex extensions exist. In regression analysis, logistic regression[1] (or logit regression) is estimating the parameters of a logistic model (a form of binary regression). Mathematically, a binary logistic model has a dependent variable with two possible values, such as pass/fail which is represented by an indicator variable, where the two values are labeled "0" and "1"

LOGISTIC REGRESSION

```
In [33]: # LogisticRegression
LG = LogisticRegression(C=1, max_iter = 3000)

LG.fit(x_train, y_train)

y_pred_train = LG.predict(x_train)
print('Training accuracy is {}'.format(accuracy_score(y_train, y_pred_train)))
y_pred_test = LG.predict(x_test)
print('Test accuracy is {}'.format(accuracy_score(y_test, y_pred_test)))
print(confusion_matrix(y_test, y_pred_test))
print(classification_report(y_test, y_pred_test))
```

Training accuracy is 0.9595520103134316

Test accuracy is 0.9552974598930482

[[42729 221]

[1919 3003]]

	precision	recall	f1-score	support
0	0.96	0.99	0.98	42950
1	0.93	0.61	0.74	4922
accuracy			0.96	47872
macro avg	0.94	0.80	0.86	47872
weighted avg	0.95	0.96	0.95	47872

DECISION TREE CLASSIFIER: Decision trees use **multiple algorithms to decide to split a node into two or more sub-nodes**. The creation of sub-nodes increases the homogeneity of resultant sub-nodes. ... The decision tree splits the nodes on all available variables and then selects the split which results in most homogeneous sub-nodes.

DECISION TREE CLASSIFIER

```
n [34]: # DecisionTreeClassifier
DT = DecisionTreeClassifier()

DT.fit(x_train, y_train)
y_pred_train = DT.predict(x_train)
print('Training accuracy is {}'.format(accuracy_score(y_train, y_pred_train)))
y_pred_test = DT.predict(x_test)
print('Test accuracy is {}'.format(accuracy_score(y_test, y_pred_test)))
print(confusion_matrix(y_test, y_pred_test))
print(classification_report(y_test, y_pred_test))

Training accuracy is 0.9988898736783678
Test accuracy is 0.9385444518716578
[[41557 1393]
 [ 1549  3373]]
              precision    recall  f1-score   support

         0       0.96      0.97      0.97       42950
         1       0.71      0.69      0.70        4922

 accuracy          0.94          0.94          0.94       47872
  macro avg       0.84          0.83          0.83       47872
 weighted avg     0.94          0.94          0.94       47872
```

RANDOME FOREST CLASSIFIER: Random forest, like its name implies, consists of a large number of individual decision trees that operate as an ensemble. Each individual tree in the random forest spits out a class prediction and the class with the most votes becomes our model's prediction. The reason for this wonderful effect is that the trees protect each other from their individual errors

RANDOM FOREST CLASSIFIER

```
Out[67]: {'colsample_bytree': 0.7,
          'learning_rate': 0.07,
          'max_depth': 5,
          'min_child_weight': 4,
          'n_estimators': 500,
          'nthread': 4,
          'objective': 'reg:linear',
          'silent': 1,
          'subsample': 0.7}
```

XGBosst Classifier:

XGBoost is a **decision-tree-based ensemble Machine Learning algorithm** that uses a gradient boosting framework. In prediction problems involving unstructured data (images, text, etc.) ... A wide range of applications: Can be used to solve regression, classification, ranking, and user-defined prediction problems.

XGBoost

```
In [36]: # xgboost
import xgboost
xgb = xgboost.XGBClassifier()
xgb.fit(x_train, y_train)
y_pred_train = xgb.predict(x_train)
print('Training accuracy is {}'.format(accuracy_score(y_train, y_pred_train)))
y_pred_test = xgb.predict(x_test)
print('Test accuracy is {}'.format(accuracy_score(y_test, y_pred_test)))
print(confusion_matrix(y_test, y_pred_test))
print(classification_report(y_test, y_pred_test))
```

C:\Users\sarmi\anaconda3\lib\site-packages\xgboost\sklearn.py:888: UserWarning: The use of label encoder in XGBClassifier and will be removed in a future release. To remove this warning, do the following: 1) Pass option use_label_encoder=False when initializing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_classes-1]. See https://xgboost.ai/docs/#faq-frequent-warnings for more details. warnings.warn(label_encoder_deprecation_msg, UserWarning)

[20:24:54] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.3.0/src/learner.cc:1061: Starting in fault evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicit you'd like to restore the old behavior.

Training accuracy is 0.9614052050600274
Test accuracy is 0.9526236631016043

```
[[42689 261]
 [2007 2915]]
```

	precision	recall	f1-score	support
0	0.96	0.99	0.97	42950
1	0.92	0.59	0.72	4922
accuracy			0.95	47872
macro avg	0.94	0.79	0.85	47872
weighted avg	0.95	0.95	0.95	47872

ADABOOST: An AdaBoost [1] classifier is a meta-estimator that begins by fitting a classifier on the original dataset and then fits additional copies of the classifier on the same dataset but where the weights of incorrectly classified instances are adjusted such that subsequent classifiers focus more on difficult cases.

Adaboost classifier

```
In [37]: #AdaBoostClassifier
ada=AdaBoostClassifier(n_estimators=100)
ada.fit(x_train, y_train)
y_pred_train = ada.predict(x_train)
print('Training accuracy is {}'.format(accuracy_score(y_train, y_pred_train)))
y_pred_test = ada.predict(x_test)
print('Test accuracy is {}'.format(accuracy_score(y_test,y_pred_test)))
print(confusion_matrix(y_test,y_pred_test))
print(classification_report(y_test,y_pred_test))
```

```
Training accuracy is 0.951118631321677
Test accuracy is 0.9490307486631016
[[42553   397]
 [ 2043  2879]]

              precision    recall  f1-score   support

     0           0.95        0.99        0.97        42950
     1           0.88        0.58        0.70         4922

 accuracy                   0.95        47872
 macro avg           0.92        0.79        0.84        47872
 weighted avg           0.95        0.95        0.94        47872
```

K-Neighbors classifier: The k-nearest neighbors (KNN) algorithm is a simple, supervised machine learning algorithm that can be used to solve both classification and regression problems. It's easy to implement and understand, but has a major drawback of becoming significantly slower as the size of that data in use grows.

K-Neighbors classifier

```
In [38]: #KNeighborsClassifier
knn=KNeighborsClassifier(n_neighbors=9)
knn.fit(x_train, y_train)
y_pred_train = knn.predict(x_train)
print('Training accuracy is {}'.format(accuracy_score(y_train, y_pred_train)))
y_pred_test = knn.predict(x_test)
print('Test accuracy is {}'.format(accuracy_score(y_test,y_pred_test)))
print(confusion_matrix(y_test,y_pred_test))
print(classification_report(y_test,y_pred_test))
```

```
Training accuracy is 0.922300110117369
Test accuracy is 0.9173629679144385
[[42809   141]
 [ 3815  1107]]

              precision    recall  f1-score   support

     0           0.92        1.00        0.96        42950
     1           0.89        0.22        0.36         4922

 accuracy                   0.92        47872
 macro avg           0.90        0.61        0.66        47872
 weighted avg           0.91        0.92        0.89        47872
```

- Key Metrics for success in solving problem under consideration

We found the accuracy, confusion matrix and then decided that Random forest is the best model. Tuning the best model further.

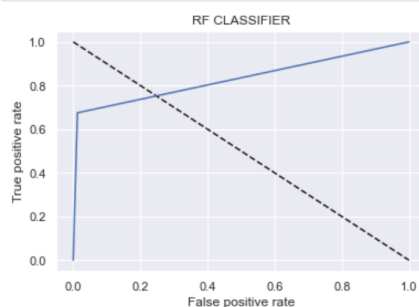
Random Forest Classifier

```
In [39]: # RandomForestClassifier
RF = RandomForestClassifier()
RF.fit(x_train, y_train)
y_pred_train = RF.predict(x_train)
print('Training accuracy is {}'.format(accuracy_score(y_train, y_pred_train)))
y_pred_test = RF.predict(x_test)
print('Test accuracy is {}'.format(accuracy_score(y_test, y_pred_test)))
cvs=cross_val_score(RF, x, y, cv=10, scoring='accuracy').mean()
print('cross validation score :', cvs*100)
print(confusion_matrix(y_test, y_pred_test))
print(classification_report(y_test, y_pred_test))
```

```
Training accuracy is 0.9988898736783678
Test accuracy is 0.9554019050802139
cross validation score : 95.67026523523793
[[42414  536]
 [ 1599  3323]]
```

		precision	recall	f1-score	support
	0	0.96	0.99	0.98	42950
	1	0.86	0.68	0.76	4922
	accuracy			0.96	47872
	macro avg	0.91	0.83	0.87	47872
	weighted avg	0.95	0.96	0.95	47872

```
In [40]: #Plotting the graph which tells us about the area under curve , more the area under curve more will be the better prediction
# model is performing good :
fpr, tpr, thresholds=roc_curve(y_test, y_pred_test)
roc_auc=auc(fpr, tpr)
plt.plot([0,1],[1,0], 'k--')
plt.plot(fpr, tpr, label = 'RF Classifier')
plt.xlabel('False positive rate')
plt.ylabel('True positive rate')
plt.title('RF CLASSIFIER')
plt.show()
```



CONCLUSION

- Conclusions

According to the performance metrics, Random forest classifier is the best among the others.

• Limitations & Scope for Future

Although we have tried quite several parameters in refining my model, there can exist a better model which gives greater accuracy.

In future, we may be able to add the following to the existing model for enhanced performance

The current project predicts the type or toxicity in the comment.

- Analyse which age group is being toxic towards a particular group or brand.
- Add feature to automatically sensitize words which are classified as toxic.
- Automatically send alerts to the concerned authority if threats are classified as severe.
- Build a feedback loop to further increase the efficiency of the model.
- Handle mistakes and short forms of words to get better accuracy of the result